


Chapter 5: CPU Scheduling





Chapter 5: CPU Scheduling

- ❑ Basic Concepts
- ❑ Scheduling Criteria
- ❑ Scheduling Algorithms
- ❑ Thread Scheduling
- ❑ Multiple-Processor Scheduling
- ❑ Real-Time CPU Scheduling
- ❑ Operating Systems Examples
- ❑ Algorithm Evaluation





Objectives

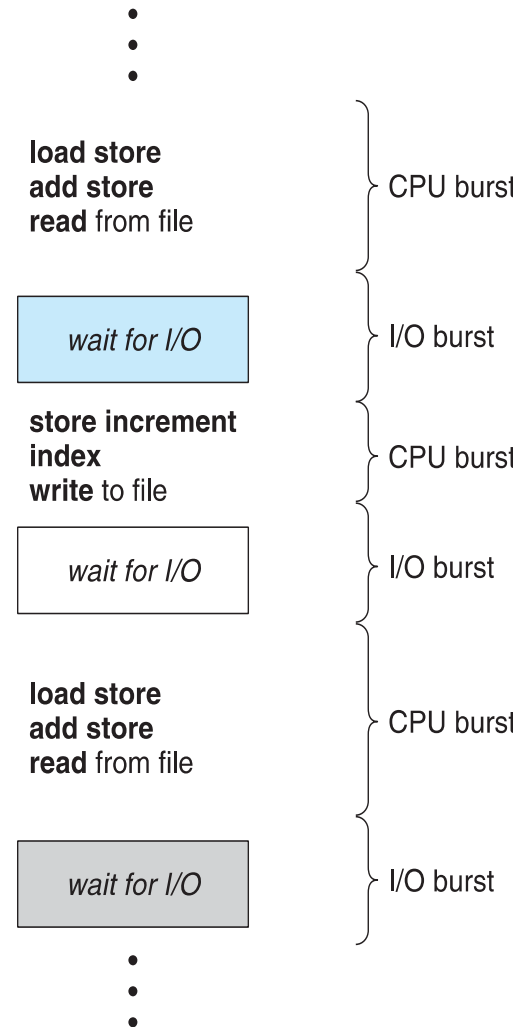
- ❑ To introduce CPU scheduling, which is the basis for multiprogrammed operating systems
- ❑ To describe various CPU-scheduling algorithms
- ❑ To discuss evaluation criteria for selecting a CPU-scheduling algorithm for a particular system
- ❑ To examine the scheduling algorithms of several operating systems





Basic Concepts

- ❓ Maximum CPU utilization obtained with multiprogramming
- ❓ CPU–I/O Burst Cycle – Process execution consists of a **cycle** of CPU execution and I/O wait
- ❓ **CPU burst** followed by **I/O burst**
- ❓ CPU burst distribution is of main concern





CPU Scheduler

- ❑ **Short-term scheduler** selects from among the processes in ready queue, and allocates the CPU to one of them
 - ❑ Queue may be ordered in various ways
- ❑ CPU scheduling decisions may take place when a process:
 1. Switches from running to waiting state
 2. Switches from running to ready state
 3. Switches from waiting to ready
 4. Terminates
- ❑ Scheduling under 1 and 4 is **nonpreemptive**
- ❑ All other scheduling is **preemptive**
 - ❑ Consider access to shared data
 - ❑ Consider preemption while in kernel mode
 - ❑ Consider interrupts occurring during crucial OS activities





Dispatcher

- ❑ Dispatcher module gives control of the CPU to the process selected by the short-term scheduler; this involves:
 - ❑ switching context
 - ❑ switching to user mode
 - ❑ jumping to the proper location in the user program to restart that program
- ❑ **Dispatch latency** – time it takes for the dispatcher to stop one process and start another running





Scheduling Criteria

- ❑ **CPU utilization** – keep the CPU as busy as possible
- ❑ **Throughput** – # of processes that complete their execution per time unit
- ❑ **Turnaround time** – amount of time to execute a particular process
- ❑ **Waiting time** – amount of time a process has been waiting in the ready queue
- ❑ **Response time** – amount of time it takes from when a request was submitted until the first response is produced, not output (for time-sharing environment)





Scheduling Algorithm Optimization Criteria

- ❑ Max CPU utilization
- ❑ Max throughput
- ❑ Min turnaround time
- ❑ Min waiting time
- ❑ Min response time





First- Come, First-Served (FCFS) Scheduling

<u>Process</u>	<u>Burst Time</u>
P_1	24
P_2	3
P_3	3

- ❓ Suppose that the processes arrive in the order: P_1, P_2, P_3
The Gantt Chart for the schedule is:



- ❓ Waiting time for $P_1 = 0$; $P_2 = 24$; $P_3 = 27$
❓ Average waiting time: $(0 + 24 + 27)/3 = 17$





FCFS Scheduling (Cont.)

Suppose that the processes arrive in the order:

$$P_2, P_3, P_1$$

[?] The Gantt chart for the schedule is:



[?] Waiting time for $P_1 = 6$; $P_2 = 0$; $P_3 = 3$

[?] Average waiting time: $(6 + 0 + 3)/3 = 3$

[?] Much better than previous case

[?] **Convoy effect** - short process behind long process

[?] Consider one CPU-bound and many I/O-bound processes





Round Robin (RR)

- ❑ Each process gets a small unit of CPU time (**time quantum** q), usually 10-100 milliseconds. After this time has elapsed, the process is preempted and added to the end of the ready queue.
- ❑ If there are n processes in the ready queue and the time quantum is q , then each process gets $1/n$ of the CPU time in chunks of at most q time units at once. No process waits more than $(n-1)q$ time units.
- ❑ Timer interrupts every quantum to schedule next process
- ❑ Performance
 - ❑ q large \Rightarrow FIFO
 - ❑ q small $\Rightarrow q$ must be large with respect to context switch, otherwise overhead is too high









Example of RR with Time Quantum = 4

<u>Process</u>	<u>Burst Time</u>
P_1	24
P_2	3
P_3	3

 The Gantt chart is:



-  Typically, higher average turnaround than SJF, but better **response**
-  q should be large compared to context switch time
-  q usually 10ms to 100ms, context switch < 10 usec





Chapter 5: CPU Scheduling

- [?] Basic Concepts**
- [?] Scheduling Criteria**
- [?] Scheduling Algorithms**
- [?] Thread Scheduling**
- [?] Multiple-Processor Scheduling**
- [?] Real-Time CPU Scheduling**
- [?] Operating Systems Examples**
- [?] Algorithm Evaluation**

Quiz:

- 1. What is the main disadvantage of the First-Come, First-Served (FCFS) scheduling algorithm?**
- 2. What does the Round Robin (RR) algorithm do when a process uses up its time quantum?**
- 3. What is non-preemptive scheduling?**
- 4. What is preemptive scheduling?**

