


# Chapter 5: CPU Scheduling

---





# Multiple-Processor Scheduling

---

- ☐ Multiple processors?
- ☐ CPU scheduling more complex when multiple CPUs are available
- ☐ **Asymmetric multiprocessing** – only one processor accesses the system data structures, alleviating the need for data sharing
- ☐ **Symmetric multiprocessing (SMP)** – each processor is self-scheduling, all processes in common ready queue, or each has its own private queue of ready processes
  - ☐ Currently, most common





## Multiple-Processor Scheduling – Load Balancing

---


- If SMP, need to keep all CPUs loaded for efficiency
- **Load balancing** attempts to keep workload evenly distributed
  - **Push migration** – periodic task checks load on each processor, and if found pushes task from overloaded CPU to other CPUs
  - **Pull migration** – idle processors pulls waiting task from busy processor





# Priority-based Scheduling

- For real-time scheduling, scheduler must support preemptive, priority-based scheduling
  - But only guarantees soft real-time
- For hard real-time must also provide ability to meet deadlines
- Processes have new characteristics: **periodic** ones require CPU at constant intervals
  - Has processing time  $t$ , deadline  $d$ , period  $p$
  - $0 \leq t \leq d \leq p$
  - **Rate** of periodic task is  $1/p$





# Rate Monotonic Scheduling

- A priority is assigned based on the inverse of its period
- Shorter periods = higher priority;
- Longer periods = lower priority
- Example with three processors

Process	Capacity	Period
P1	3	20
P2	2	5
P3	2	10





## Earliest Deadline First Scheduling (EDF)

- Priorities are assigned according to deadlines:

the earlier the deadline, the higher the priority;

the later the deadline, the lower the priority

**Process Capacity Period deadline**

<b>P1</b>	<b>3</b>	<b>20</b>	<b>7</b>
<b>P2</b>	<b>2</b>	<b>5</b>	<b>4</b>
<b>P3</b>	<b>2</b>	<b>10</b>	<b>8</b>





## Activity

### Apply Rate Monotonic Scheduling

Process	Capacity	Period
P1	5	15
P2	4	5
P3	4	3

Process	Capacity	Period
P1	1	15
P2	3	5
P3	1	3


EDF:

Process	Capacity	Period	Deadline
P1	1	4	4
P2	2	6	6
P3	3	8	8



# End of Chapter 5


---





# Chapter 6: Synchronization Tools

---





# Chapter 6: Synchronization Tools

---

- Background
- The Critical-Section Problem
- Peterson's Solution
- Synchronization Hardware
- Mutex Locks
- Semaphores





# Objectives

---

To present the concept of process synchronization.

To introduce the critical-section problem, whose solutions can be used to ensure the consistency of shared data

To present both software and hardware solutions of the critical-section problem

To examine several classical process-synchronization problems

To explore several tools that are used to solve process synchronization problems





# Background

---

Cooperating process - Processes that can affect or get affected by other processes executing in the system.

Cooperating processes can be:

- Directly sharing a logical address space(both code and data)

- Allowed to share data only through files or messages

Problem: Concurrent access to shared data may result in data inconsistency

Solution: Process synchronization





# Background

---

Concurrent access to shared data may result in data inconsistency

Maintaining data consistency requires mechanisms to ensure the orderly execution of cooperating processes

Eg: Producer consumer problem

Solution : with the use of shared memory

Buffer: the region in the shared memory that both producer and consumer use

The producer produce one item while the consumer is consuming another item

The producer and consumer should be synchronized





# Background

---

Buffer: the region in the shared memory that both producer and consumer use

Two types of buffer:

Unbounded — No limit on the size of the buffer. The producer can always produce items, consumer may need to wait for new items

Bounded — Fixed buffer size. Consumer must wait when buffer is empty. Producer must wait when the buffer is full

Illustration of the problem:

Suppose that we wanted to provide a solution to the consumer-producer problem that fills **all** the buffers. We can do so by having an integer **counter** that keeps track of the number of full buffers. Initially, **counter** is set to 0. It is incremented by the producer after it produces a new buffer and is decremented by the consumer after it consumes a buffer.





# Producer

---

```
while (true) {  
    /* produce an item in next produced */  
  
    while (counter == BUFFER_SIZE) ;  
        /* do nothing */  
    buffer[in] = next_produced;  
    in = (in + 1) % BUFFER_SIZE;  
    counter++;  
}
```





# Consumer

---

```
while (true) {  
    while (counter == 0)  
        ; /* do nothing */  
    next_consumed = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
    counter--;  
    /* consume the item in next consumed */  
}
```







# Race Condition

**counter++** could be implemented as

```
register1 = counter
register1 = register1 + 1
counter = register1
```

**counter--** could be implemented as

```
register2 = counter
register2 = register2 - 1
counter = register2
```

Consider this execution interleaving with “count = 5” initially:

S0: producer execute	<b>register1 = counter</b>	{register1 = 5}
S1: producer execute	<b>register1 = register1 + 1</b>	{register1 = 6}
S2: consumer execute	<b>register2 = counter</b>	{register2 = 5}
S3: consumer execute	<b>register2 = register2 - 1</b>	{register2 = 4}
S4: producer execute	<b>counter = register1</b>	{counter = 6}
S5: consumer execute	<b>counter = register2</b>	{counter = 4}

When several processes access and manipulate the data concurrently and the outcome of the execution depends on the particular order in which the access takes place — Race Condition

