


Chapter 6: Synchronization Tools





Chapter 6: Synchronization Tools

- Background
- The Critical-Section Problem
- Peterson's Solution
- Synchronization Hardware
- Mutex Locks
- Semaphores





Objectives

To present the concept of process synchronization.

To introduce the critical-section problem, whose solutions can be used to ensure the consistency of shared data

To present both software and hardware solutions of the critical-section problem

To examine several classical process-synchronization problems

To explore several tools that are used to solve process synchronization problems





Background

Cooperating process - Processes that can affect or get affected by other processes executing in the system.

Cooperating processes can be:

- Directly sharing a logical address space(both code and data)

- Allowed to share data only through files or messages

Problem: Concurrent access to shared data may result in data inconsistency

Solution: Process synchronization





Background

Concurrent access to shared data may result in data inconsistency

Maintaining data consistency requires mechanisms to ensure the orderly execution of cooperating processes

Eg: Producer consumer problem

Solution : with the use of shared memory

Buffer: the region in the shared memory that both producer and consumer use

The producer produce one item while the consumer is consuming another item

The producer and consumer should be synchronized





Background

Buffer: the region in the shared memory that both producer and consumer use

Two types of buffer:

Unbounded — No limit on the size of the buffer. The producer can always produce items, consumer may need to wait for new items

Bounded — Fixed buffer size. Consumer must wait when buffer is empty. Producer must wait when the buffer is full

Illustration of the problem:

Suppose that we wanted to provide a solution to the consumer-producer problem that fills **all** the buffers. We can do so by having an integer **counter** that keeps track of the number of full buffers. Initially, **counter** is set to 0. It is incremented by the producer after it produces a new buffer and is decremented by the consumer after it consumes a buffer.





Race Condition

counter++ could be implemented as

```
register1 = counter
register1 = register1 + 1
counter = register1
```

counter-- could be implemented as

```
register2 = counter
register2 = register2 - 1
counter = register2
```

Consider this execution interleaving with “count = 5” initially:

S0: producer execute	register1 = counter	{register1 = 5}
S1: producer execute	register1 = register1 + 1	{register1 = 6}
S2: consumer execute	register2 = counter	{register2 = 5}
S3: consumer execute	register2 = register2 - 1	{register2 = 4}
S4: producer execute	counter = register1	{counter = 6}
S5: consumer execute	counter = register2	{counter = 4}

When several processes access and manipulate the data concurrently and the outcome of the execution depends on the particular order in which the access takes place — Race Condition





Producer

```
while (true) {  
    /* produce an item in next produced */  
  
    while (counter == BUFFER_SIZE) ;  
        /* do nothing */  
    buffer[in] = next_produced;  
    in = (in + 1) % BUFFER_SIZE;  
    counter++;  
}
```





Consumer

```
while (true) {  
    while (counter == 0)  
        ; /* do nothing */  
    next_consumed = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
    counter--;  
    /* consume the item in next consumed */  
}
```





Critical Section Problem

Consider system of n processes $\{p_0, p_1, \dots p_{n-1}\}$

Each process has **critical section** segment of code

Process may be changing common variables, updating table, writing file, etc

When one process in critical section, no other may be in its critical section

Critical section problem is to design protocol to solve this

Each process must ask permission to enter critical section in **entry section**, may follow critical section with **exit section**, then **remainder section**





Critical Section

General structure of process P_i

```
do {  
    entry section  
    critical section  
    exit section  
    remainder section  
} while (true);
```





Solution to Critical-Section Problem

1. **Mutual Exclusion** - If process P_i is executing in its critical section, then no other processes can be executing in their critical sections
2. **Progress** - If no process is executing in its critical section and there exist some processes that wish to enter their critical section, then the selection of the processes that will enter the critical section next cannot be postponed indefinitely
3. **Bounded Waiting** - A bound must exist on the number of times that other processes are allowed to enter their critical sections after a process has made a request to enter its critical section and before that request is granted

https://www.youtube.com/watch?v=6x_XMDCMyAk.





Peterson's Solution

Good algorithmic description of solving the problem

Two process solution — Peterson's solution is restricted to two processes that alternate the execution between the critical section and remainder section.

The two processes share two variables:

```
int turn;  
Boolean flag[2]
```

The variable `turn` indicates whose turn it is to enter the critical section

The `flag` array is used to indicate if a process is ready to enter the critical section. `flag[i] = true` implies that process P_i is ready!





Algorithm for Process P_i

do {

```
flag[i] = true;
```

```
turn = j;
```

```
while (flag[j] && turn == j);
```

critical section

```
flag[i] = false;
```

remainder section

```
} while (true);
```





Peterson's Solution (Cont.)

Provable that the three CS requirement are met:

1. Mutual exclusion is preserved

P_i enters CS only if:

either `flag[j] = false` or `turn = i`

2. Progress requirement is satisfied
3. Bounded-waiting requirement is met





Synchronization Hardware

Many systems provide hardware support for implementing the critical section code.

All solutions below based on idea of **locking**

Protecting critical regions via locks

Uniprocessors – could disable interrupts

Currently running code would execute without preemption





Solution to Critical-section Problem Using Locks

```
do {  
    acquire lock  
        critical section  
    release lock  
        remainder section  
} while (TRUE);
```



Quiz

1. What are the two types of processes?
2. Which type of processes are in need of synchronization methods?
3. Why do we need process synchronization?
4. What are the three issues in producer consumer problem?
5. What are the two types of buffer?
6. What is a critical section?
7. What are the 4 sections in a process's structure?
8. What are the three criteria to avoid critical section problem?
9. What are the counter operations in producer consumer problem?